

AuraScribe

Product Overview & UI Reference

FOR REDESIGN — AS-BUILT STATE, v1.0.0

What this document is

A faithful description of AuraScribe as it exists today — every screen, the navigation, the layout, the interaction states, and the technical shape underneath. It is written so a designer who has never seen the app can redesign it without guessing what each surface does. It records the current visual system honestly (including where it is deliberately austere) and marks clearly what is free to change versus what must not.

1 What AuraScribe is

AuraScribe is a free, open-source, **local-first** voice dictation app for Windows. The user presses a global hotkey, speaks, and clean punctuated text is typed at their cursor in whatever application has focus — a browser, an editor, a chat box, anything that accepts typing.

Everything runs on the user's own machine. The speech model (OpenAI's Whisper, running natively) transcribes locally; a lightweight local pass cleans the text. The only network request the app ever makes is the one-time model download.

Product principles (these constrain any redesign)

- **Local-first.** Audio and text never leave the machine. No cloud, no telemetry, no analytics.
- **Free forever.** No tiers, no word caps, no account.
- **Lightweight.** ~4.4 MB installer, ~40 MB idle RAM. Every dependency is weighed.
- **Honest.** The UI never claims more than the code does.
- **Background utility.** It lives in the system tray; the window is opened when needed, not stared at.

Design implication: the product is an instrument you reach for, not a dashboard you live in. Speed and clarity beat richness. A redesign can change the whole visual language, but it should keep the app quiet, fast, and out of the way.

2 Information architecture

The app is a single window with a fixed left sidebar and a content area. The sidebar holds six destinations plus a live status rail. Two surfaces live outside the main window: a floating recording pill and the system-tray icon.

Surface	Type	Purpose
Dictate	Main tab	Home. Big mic button, hotkey hint, live meter, last transcript.
History	Tab	Every past dictation, newest first. Copy or clear.
Words	Tab	Personal dictionary: “heard as” → “written as” corrections.
Snippets	Tab	Text expansions: a trigger phrase inserts a saved block.
Insights	Tab	Usage stats: words, words/min, dictations, streak, time saved.
Settings	Tab	Voice model, hotkey, mic, cleanup, appearance, sound, startup.
Recording pill	Floating	Small always-on-top “Listening / Processing” indicator, bottom-centre.
Tray icon	System tray	Idle / listening / processing colour; left-click opens, right-click quits.

3 Screen by screen

Dictate (home)

The default screen and the one users see most. Its job is to answer one question at a glance: is the app ready, and is it hearing me?

Elements

- The **signal meter** — the signature element, and the only thing that animates (see §5).
- A large circular **mic button** to start/stop without the hotkey.
- The hotkey hint (e.g. `Ctrl Shift Space`) and the active model id.
- **Tap vs Hold** mode chips (press-to-toggle, or push-to-talk).
- The last transcript, with a **Copy** action.

States

- **No model:** replaced by an “Add a voice model to begin” prompt with a button to Settings.
- **Listening:** meter ticks in red; mic button active.
- **Processing:** meter sweeps amber; button shows a spinner.
- Errors appear inline in red beneath the button.

History

A record of what was dictated, so a user can recover or reuse earlier text.

Elements

- A list of past dictations, newest first, each with its text and a **Copy** action.
- A **Clear all** action in the header.
- Empty state: “Nothing dictated yet.”

Words (personal dictionary)

Fixes names and terms Whisper mis-hears — applied automatically to every future dictation.

Elements

- An add row: **Heard as** (e.g. “kubernetes”) → **Written as** (e.g. “Kubernetes”).
- A list of saved corrections, each removable.
- Empty state invites adding the first term.

Snippets

Text expansion: say a short trigger, get a saved block of text inserted.

Elements

- An add form: **When you say** (trigger) / **Insert this** (expansion).
- A list of saved snippets, each removable.

Insights

Lightweight usage stats, framed around the value the user gets, not vanity metrics.

Elements

- Stat tiles: Words dictated, Words per minute, Dictations, Words today, Active days, Time saved.
- Numbers use tabular figures so they don't jitter as they update.
- Empty state until there is data.

Settings

One scrollable pane, grouped into sections. The most important section is the voice model.

Elements

- **Voice model:** a card per model with size, speed, accuracy, a live speed estimate for the user's machine, and a warning before download if it would be slow. Actions: Install / In use / delete.
- **Hotkey:** rebindable combination + Tap/Hold mode.
- **Microphone:** input device picker.
- **Cleanup:** toggle automatic cleanup; sub-toggle to remove filler words.
- **Application:** start at login, **play a sound on start/stop**, appearance (light/dark/system).

4 Layout & navigation

The window opens at 1480×936 (logical), centred, and shrinks to fit smaller screens. Minimum 860×560. Structure is a two-column shell:

```
+-----+
| AuraScribe | |
| ON-DEVICE | [ signal meter ] |
| | |
| ■ Dictate | ( ● ) |
| History | Press Ctrl Shift Space |
| Words | |
| Snippets | [ Tap ] [ Hold ] |
| Insights | |
| Settings | |
| | |
| ● Ready | |
| base.en | |
+-----+
```

Sidebar

Expanded 212 px, collapsible to 60 px (icon-only, with tooltips). A flat list — no groups, no accordions; six items need no hierarchy. Each row is an icon in a fixed slot plus a label; collapsing only changes width and hides labels, so icons never shift. Pinned to the bottom is the **status rail**: a coloured dot + state word (Ready / Listening / Working / Needs a model) and the loaded model id beneath.

Content area

Centred with a max width so panels don't stretch into unreadable lines. Every tab uses the same page header pattern (title + one-line description) and the same card, empty-state, and error components, so the surfaces feel like one system.

5 Visual system (as-built)

The current direction is deliberately austere — “instrument, not dashboard.” It is documented here so the redesign knows what it is replacing. **You are free to change all of it** except the accessibility floor (§7) and the product principles (§1).

Colour — colour means state, never decoration

Token	Light	Dark	Meaning
background	#FAFBFC	#0E1116	page
card	#FFFFFF	#171B22	panel
border	#E2E7EC	#242A33	hairline
primary	#0B7C8A	#3ECEDE	signal-cyan — ready / active / live
record	#E03B45	#F0555C	recording
standby	#C4780F	#E8A33D	transcribing

Three states, three colours; cyan is reserved for “live.” Dark mode is a cool, blue-shifted graphite so panels read as anodised metal rather than flat grey.

Typography

UI text is **Inter**. Machine-reported values — hotkeys, model ids, sizes, counts, durations — are set in the system **monospace** stack. That split is what makes the app read as an instrument. Numerals use tabular figures. *The redesign brief notes the current type feels cheap; this pairing and scale are exactly what is open to change.*

Shape & surface

Radius 10 px throughout — precise, not bubbly. The current build uses **no gradients, no glass, no drop shadows, no hover-lift**; borders and background do the work. The redesign explicitly wants the opposite — soft gradients, curves, more warmth — which is a legitimate change of direction, not a violation of any principle.

Signature: the signal meter

The one memorable element and the only place with real motion. **Idle**: a flat row of dim ticks. **Listening**: ticks rise and fall in red, staggered like a waveform. **Transcribing**: a single amber band sweeps across. It answers “is this thing hearing me?” without words. Any redesign should keep a single, functional, motion-bearing focal element like this rather than scattering animation across the UI.

6 Interaction & states

The dictation loop

- User presses the hotkey (default `Ctrl Shift Space`) or the mic button.
- A soft rising two-note **chime** plays; tray + meter turn red; the “Listening” pill appears bottom-centre.
- As the user speaks, audio is transcribed in **chunks** in the background; the transcript builds up live.
- User presses again (Tap mode) or releases (Hold mode). A falling chime plays; state goes amber “Processing.”
- Only the final chunk remains to transcribe, so the wait is ~1–2 s regardless of length; text is pasted at the cursor.

Global state model

Four visible states, shown consistently in the tray icon, the floating pill, the sidebar rail, and the meter: **Needs a model, Ready, Listening, Processing**. All four read from one authoritative backend value, so the surfaces never disagree.

7 Accessibility floor (must survive any redesign)

- Visible keyboard focus rings (2 px accent) on every interactive element.
- `prefers-reduced-motion` disables the meter animation and all transitions.
- State is conveyed by **text and shape as well as colour** — never colour alone.
- Toggles are real `role="switch"` controls with `aria-checked`.
- Colour contrast holds in both themes (light-mode cyan is darkened to `#0B7C8A` for contrast on white).

8 Voice models

Transcription runs on the CPU, so the list is now restricted to models that keep up with speech on an ordinary machine. The heavy `large-v3` models were removed: on a CPU they ran many times slower than real time and overheated the machine. They can return behind a GPU build later.

Model	Size	Language	Speed (CPU)	Role
tiny.en	~75 MB	English	fastest, ~0.2×	weak machines
base.en	~142 MB	English	fast, ~0.5×	recommended
base	~142 MB	Multilingual	fast, ~0.6×	other languages

“Speed” is processing time as a multiple of the speech length; below 1.0 means it finishes faster than you spoke. The app recommends the most accurate model that still keeps up.

9 Technical shape (context for designers)

You don't need this to redesign the surfaces, but it explains what is cheap and what is hard to change.

Layer	Technology	Notes
Shell	Tauri 2 + Rust	Native window, tray, global hotkey, all dictation logic.
UI	Next.js 14 + React 18 + TypeScript	Settings window and the recording pill only. Tailwind CSS.
Speech	Whisper.cpp (native)	Runs in-process on the CPU; models downloaded once.
Cleanup	Local Rust string pass	Fillers, punctuation, question marks, casing. No LLM, no network.
Storage	SQLite (local)	Settings, history, dictionary, snippets — all on-device.
Injection	Win32 SendInput / clipboard	Types or pastes the result into the focused app.

What this means for design

- The UI is standard web tech (React + Tailwind), so **visual redesign is cheap** — it is HTML/CSS.
- The recording pill is a separate tiny transparent window; it must stay lightweight and legible over any background.
- There is no server, so anything requiring cloud (AI rewriting, sync, accounts) is out of scope by principle.
- Themes are driven by CSS variables; light and dark must both be designed.

10 Redesign brief — free vs fixed

Free to change	Must not change
Typography, type scale, weights	Local-first: no cloud, no telemetry
Colour palette & gradients	Free forever: no paywalls, no accounts
Corner radius, curves, softness	Lightweight footprint (weigh new assets/fonts)
Card / panel / button styling	Accessibility floor (\$7)
Iconography and illustration	Colour still communicates state clearly
Spacing, density, motion feel	Stays a quiet background utility, fast to use
The meter's visual form	A single functional focal element, not scattered animation

Direction the owner has asked for

Warmer and more premium than the current austere “instrument” look: soft, light gradients that are easy on the eyes; rounded, friendly forms; more considered typography; an overall feel that reads as a polished consumer product rather than a developer tool — while keeping it fast, calm, and legible in both light and dark.

Generated from the AuraScribe source as-built. Screens, tokens, and behaviour reflect the current code, not an aspiration.